SUPPLEMENT: SPECULATIVE DECODING ACROSS CONSUMER SYSTEMS

Anonymous Authors[1]

A SUPPLEMENT SCOPE

This separate appendix accompanies the measured hardware and runtime studies in manuscript v0.7. It supplies statistical, measurement-boundary, and provenance details. Internal source paths, repository identifiers and evidence hashes remain in the private evidence ledger rather than the anonymous PDF.

B FROZEN ESTIMATORS

Table A1 preserves each campaign's estimator. The primary sample unit is the prompt. Repeated observations contribute to its summary, rather than being treated as independent prompts. Pair ratios divide matched baseline decode time by speculative decode time. Historical M4, RTX, and B570 equality counts compare re-tokenized parsed CLI output. In the historical M4 HumanEval pipeline the parsed text includes the prompt; equality of those arrays is not a direct generated-ID comparison. B570 strips the prompt and normalizes the text before re-tokenization, which still does not recover a direct-ID trace. The 358H runner and prospective M4 study record generated IDs directly. A geometric mean across prompt summaries is used for headline speedup. B570 timing remains diagnostic regardless of its statistical estimator.

The B570 finalizer builds its prompt vector in the insertion order of the archived pair records. Sorting by task ID changes the finite fixed-seed bootstrap endpoints even though it leaves the underlying estimator unchanged. Exact replay retains that original order. The M4 HumanEval lower endpoint uses a different order statistic from M4 24Q and RTX; those frozen conventions are retained. None of these prompt intervals captures the full distribution of possible thermal, session or runtime states.

C ENERGY BOUNDARIES AND ALLOWED CLAIMS

Ratios above one favor speculation within the named boundary. For M4 fixed-$k = 2$, the package ratio is 1.013 [0.990, 1.034], so a package-energy improvement is not established. The registered $p_{\min} = 0.9$ secondary condition has a package-efficiency ratio 1.041 [1.018, 1.065]. This is not a replacement for the primary condition and does not establish a material additional speed gain. B570 serialized board efficiency is 0.302 [0.285, 0.320]; correctness failure restricts it to diagnosis. No between-platform energy ranking is valid.

D ROUND ACCOUNTING AND STATE SEMANTICS

The 358H target keeps its low-level context across rounds. Candidate evaluation, invalid-suffix rollback when needed, and correction or bonus decoding belong to the complete target-side timing region. Every OpenVINO proposal call receives the entire committed prefix. Its synchronous timing combines reconstruction and proposal; neither term can be isolated from this experiment without changing the measured execution.

The CPU path records 11,520 committed output tokens over 4,677 rounds; the NPU path records 11,520 over 4,686 rounds. Their directly measured C values are 2.4631 and 2.4584. Using the arithmetic baseline time per output token and arithmetic component totals yields CPU $(D, R, X) = (0.4304, 2.9767, 0.00124)$ and NPU $(4.1251, 2.9435, 0.00130)$. The resulting ratios 0.7227 and 0.3477 match ratios of total times, not the prompt-GM headline estimator.

Historical M4/RTX accepted-draft counters exclude correction or bonus output and have output-boundary effects. They are not substituted for the directly recorded 358H committed-progress metric. Likewise, target-forward

[1]Anonymous Institution, Anonymous City, Anonymous Region, Anonymous Country. Correspondence to: Anonymous Author <anon.email@domain.com>.

Preliminary work. Under review by the Machine Learning and Systems (MLSys) Conference. Do not distribute.

Table A1. Frozen statistical policies. GM = geometric mean. Indices below are zero-based order statistics of sorted bootstrap draws; NumPy quantiles use the original analyzer's convention.

Campaign	Within prompt	Resamples	Seed	CI endpoints / generator
M4 24Q	GM of 3 ratios	10,000	20260907	[249,9749]; Python Random
M4 HumanEval	GM of 3 ratios	100,000	20260910	[2500,97499]; Python Random
RTX HumanEval	GM of 3 ratios	10,000	20260910	[249,9749]; Python Random
B570 fallback	GM of 3 ratios	10,000	20260910	[250,9750]; Python Random
358H Phase 3	Median of 3 ratios	10,000	20260919	Quantiles .025/.975; NumPy default_rng

Table A2. Sensor boundaries and allowed interpretations. No cells are combined into whole-system energy.

Campaign	Measurement boundary	Allowed statement
M4 HumanEval	Silemetry CPU + GPU + ANE package aggregate; separate DRAM excluded	Fixed-$k = 2$ package-energy interval includes parity; registered threshold variant secondary
RTX 5070 Ti Laptop	NVIDIA board / CPU package are distinct domains; follow-up process window includes loading	No formal serialized HumanEval energy comparison
B570 fallback	Arc adapter board only; no CPU package or wall power	Diagnostic board efficiency only; path not globally exact
358H Phase 3	No qualified end-to-end energy boundary	No end-to-end CPU claim; NPU energy not admissible

PP2/PP1 ratios are never presented as complete realized target-round cost.

E Deviations and Provenance Boundaries

The M4 corrected synthetic ceiling was preregistered but not executed. Earlier synthetic acceptance diagnostics are not a substitute. The representative and separately preregistered code workloads stay distinct. Early MLX precision/correctness failures are qualification history, not failures of the final Metal Q8_0 path.

The B570 formal v2 correctness failure triggered its stop rule. A later six-prompt serialized probe passed 18/18 text-path comparisons, but full-workload serialization yielded only 60/90 under the same comparator. Earlier remote-interaction-contaminated verification observations are excluded from the plotted final serialized cells. A legacy "exact" field label in historical artifacts does not establish workload-global equality.

Windows follow-up derived tables and numerical captures support the manuscript, but historical raw per-run provenance is weaker. The CUDA score capture came from a separate same-commit instrumented build and does not replace formal measurements. The 358H source reconstruction remains `STRONGLY_SUPPORTED`, not proven exact historical upstream identity. Its completed Phase-3 records supersede the blocked-preflight empty scaffold. These are retained qualifications, not unresolved invitations to alter frozen experiments.

F Verification Context Lookup

Table A3 retains all historical context cells behind Panel A of the main verification figure. Values are target-forward ratios, not complete target-side round cost or integrated speedup. Context ranges do not represent uncertainty intervals. RTX denotes RTX 5070 Ti Laptop throughout this appendix.

G Evidence Inventory

Table A4 collects the observed regimes and evidence classes. These are case descriptions, not a hardware ranking, population-wide failure rate, or exhaustive taxonomy. M4's separate-session studies establish opposite observed outcomes; neither diagnosed mechanism is assigned to that contrast.

Table A3. Verification-path lookup. Each row is a measured context cell; ratios are shown to six decimal places from the frozen source. Qualification and equality labels refer to the historical text-path comparator for the associated execution path; they do not certify generated-ID equality. M4 contexts and B570 contexts are prefix lengths; RTX uses the reported context depth. The stacks and context sets are not matched. The original rounded ranges are M4 Metal 1.02–1.08, RTX batched 0.98–1.01, RTX serialized 2.14–2.43, B570 fast 1.24–1.64, and B570 serialized 1.98–1.99; these span contexts, not confidence intervals.

Path	Context (tokens)	V_2^{micro}	Qualification
M4 Metal	512	1.022182	Qualified shared path
M4 Metal	2048	1.076126	Qualified shared path
M4 Metal	8192	1.032264	Qualified shared path
RTX batched	0	1.009030	Correctness-blocked; diagnostic
RTX batched	256	0.983148	Correctness-blocked; diagnostic
RTX serialized	0	2.430996	Exact secondary control
RTX serialized	256	2.137011	Exact secondary control
B570 fast	512	1.242576	45/90 exact; diagnostic, inadmissible
B570 fast	2048	1.637295	45/90 exact; diagnostic, inadmissible
B570 serialized	512	1.982668	60/90 exact; diagnostic, inadmissible
B570 serialized	2048	1.991707	60/90 exact; diagnostic, inadmissible

Table A4. Evidence inventory across tested regimes. P: primary; S: secondary; M: mechanism; D: diagnostic. Both B570 paths remain globally inadmissible diagnostics. RTX serialized controls group CPU-to-RTX and RTX-to-RTX, each 90/90 text-path-equal. The 358H P/M entries separate primary outcomes from mechanism-supported decomposition. Historical M4/RTX/B570 equality labels in the frozen inventory denote re-tokenized-text agreement; 358H denotes direct-ID agreement. Neither establishes functional correctness. Case annotations are descriptive and non-exhaustive.

Case	Qualification	Outcome	Class	Case annotation
M4 representative 24Q	Qualified; 72/72 exact	Negative	P	Workload-conditioned
M4 HumanEval30	Qualified; 90/90 exact	Profitable	P	Workload-conditioned
RTX batched	Blocked; gate failed	Not admissible	M	Verification-constrained
RTX serialized controls	Qualified; 90/90 each exact	Negative	S	Verification-constrained
B570 fast	Inadmissible; 45/90 exact	Not admissible	D	Correctness-blocked
B570 serialized	Inadmissible; 60/90 exact	Not admissible	D	Correctness-blocked
358H CPU	Qualified; 90/90 exact	Negative	P/M	Target-side dominated 87.34% target side
358H NPU	Qualified; 90/90 exact	Negative	P/M	State-interface bound 58.35% replay + proposal

H RUNTIME STUDY REPRODUCIBILITY

The M4 runtime estimator first takes each condition's median of 3 repetitions, then forms a prompt latency ratio and its geometric mean across 60 prompts. The offline oracle uses the smaller target-only or Always-Spec median for each prompt. Bootstrap intervals use 10,000 draws with Python Random seed 20260924, resampling prompts in the frozen held-out order. Percentile endpoints linearly interpolate the sorted bootstrap draws at quantiles .025 and .975. A single random stream computes Target/Guard, Target/Always-Spec, and Guard/oracle intervals in that order. This preserves the frozen finite-bootstrap calculation; it does not replace the hardware estimators in Table A1. Intervals condition on the selected policy and observed prompt medians; they do not incorporate recalibration or between-session variation. The oracle selects between two measured fixed conditions, not all possible switching policies, so it is not a universal latency lower bound.

Calibration, qualification, and held-out IDs are disjoint. Selection ranks tasks with no recorded prior M4 exposure in the frozen inventory by a seeded hash within each of HumanEval and Spec-Bench; each split is balanced across these two families. The registered grid has nine candidates: first decision round $E \in \{2, 4, 8\}$ and margin $m \in \{0, 0.05, 0.10\}$. Calibration selects the exact candidate with the lowest geometric-mean latency ratio to the offline oracle. Exact ties break by mean measured guard-compute overhead, smaller exploration count, larger margin, and lexicographic candidate ID. The frozen winner is `E8_M000`; held-out results never select or retune it.

The portable M4 archive contains original observations with generated-token IDs, complete-attempt manifests, qualification and calibration records, frozen source/build information, and analysis. An independent offline verifier reproduces policy selection, qualification and formal direct-ID equality, prompt medians, ratios, intervals, and fallback fraction without executing the models. No runtime energy result is claimed.

The reviewer-requested Guard/Always-Spec contrast is post-study paired secondary analysis, not the registered primary endpoint. Aggregate results were already known; the committed analysis plan is a procedural timestamp, not preregistration or blinding. For each frozen held-out prompt i, form $r_i = \mathrm{median}(T_{\mathrm{Guard},i}) / \mathrm{median}(T_{\mathrm{Always},i})$ over its 3 repetitions, then compute $\exp(60^{-1} \sum_i \log r_i)$. Values above one mean Guard is slower. A fresh standalone Python Random stream with seed 20260924 resamples 60 paired prompts per draw in frozen held-out order, using 10,000 draws and the same linear percentile interpolation; it does not continue the primary stream. The ratio is 1.0271 [1.0130, 1.0436], with the interval above parity. Descriptive prompt counts are 24 Guard-faster, 36 Always-Spec-faster, and 0 exact ties. Both new implementations reconstruct raw medians from the same hash-verified archive and match the primary medians exactly; an independent calculation reproduces both finite-bootstrap endpoints. No new measurement or analysis choice after seeing the result was used. The interval excludes session and recalibration variation; source hashes and all prompt ratios remain in the private evidence ledger.

A metadata-recording defect requires explicit provenance recovery. Loop-variable shadowing caused the frozen helper to return the build-manifest path rather than the calibration-freeze path. Qualification and formal-start records therefore contain the build-manifest digest in the intended freeze-hash field. A versioned recovery analyzer verified the actual freeze against the clean formal-start commit, qualification, frozen source/model/calibration/policy hashes, and complete attempts before applying the original formulas. The archive retains the original markers and pipeline STOPPED state. This is a metadata-recording defect, not a measurement mutation.

The 358H diagnostic is pinned separately from the hardware campaign. Its compact evidence records reproducible target-only same-prefix divergence and proposal playback without live NPU inference, but does not establish a specific rollback-boundary bug. It explicitly records that no formal runtime performance data exist and that failed qualification prevented formal runtime performance evaluation. Internal evidence records retain the exact source commit and file hashes; the local raw diagnostic traces are not included in that compact record. No 358H runtime speedup is inferred from these diagnostics.

Table A5. Separately frozen direct-ID qualification studies. Each path must match all 90 paired sequences; exact prompts match all three repetitions. These studies do not estimate performance.

Platform	Path	Exact pairs	Exact prompts	Qualification	Comparator
RTX 5070 Ti Laptop	Batched	75/90	25/30	BLOCKED	Direct generated IDs
RTX 5070 Ti Laptop	Serialized	90/90	30/30	QUALIFIED	Direct generated IDs
B570	Fast	60/90	20/30	BLOCKED	Direct generated IDs
B570	Serialized	78/90	26/30	BLOCKED	Direct generated IDs

I Prospective Direct-ID Qualification

Table A5 reports new studies separately from the historical context lookup and timing inventory. Both freeze the same HumanEval30 IDs, Qwen3-8B and Qwen3-0.6B Q8_0 models, greedy $k = 2$, three repetitions, cap 128, and a 90/90 sequence-equality rule. Prompt IDs and uncommitted proposals are excluded; an earlier generated terminal EOG is included. No mismatch is discarded and no subset qualifies. Historical path-control outcomes informed the replications; prospective freezing does not imply blinded discovery. Each study completes 270 executions; independent offline implementations reproduce all 180 comparisons, including first divergences. Operational elapsed fields do not establish formal performance.

The RTX original embedded-server CLI exposes native committed-token debug events without a source patch. It uses context 4096, reference target micro-batch one, candidate micro-batches two/one, driver 616.92, and full CUDA0 offload. The historical RTX driver was 616.56; this is not an exact replay of historical host state. The incomplete RTX v1 transport attempt is preserved and supplies no v2 formal pair; completed v2 uses lossless byte transport, a new prospective freeze, and a complete fresh schedule.

B570 uses source instrumentation to record committed IDs in sidecars, context 1024, and a separately frozen completion reference. Its endpoint selects the first 128 committed IDs or the inclusive earlier EOG; raw CLI overshoot is retained. Vulkan0 uses driver 32.0.101.9033 with full target 37/37 and draft 29/29 layer offload. The initial Task Scheduler launch failed during CP936/UTF-8 environment-metadata decoding before inference, with zero formal executions. The first data-producing attempt directly ran the unchanged frozen runner: no semantic, code, binary, model, driver, or parameter change occurred. This deployment deviation is preserved in provenance.

The runtime/front-end, reference, and context differences preclude a causal hardware comparison. Neither study estimates population failure rates or functional code correctness. Repeat stability holds for every prompt in each condition, but stable disagreement remains disagreement. Historical B570 fast 45/90 and serialized 60/90 text-path counts and historical RTX controls retain their original comparators. New qualification does not retrospectively authorize old timing as strict-ID performance. Portable validators hash included evidence and check recorded identities of omitted external weights/binaries; they do not rehash omitted assets locally.