The shared M4 Metal path has a separately measured V_2^{micro} range of 1.022–1.076, yet the separate HumanEval30 ($k = 2$) and representative-24Q studies yield 1.387× (90/90 exact) and 0.842× (72/72 exact). RTX batched verification costs 0.983–1.009 baseline steps but fails qualification, so it has no admissible speedup. At 8,192 tokens, isolated NPU/CPU drafting reaches 48.08/30.81 tokens/s, while the qualified 358H NPU path yields 0.352× (90/90 exact). Speedups cover post-prefill decode; the detailed figures retain campaign-specific intervals and measurement boundaries. These heterogeneous observations do not isolate workload-only causality.