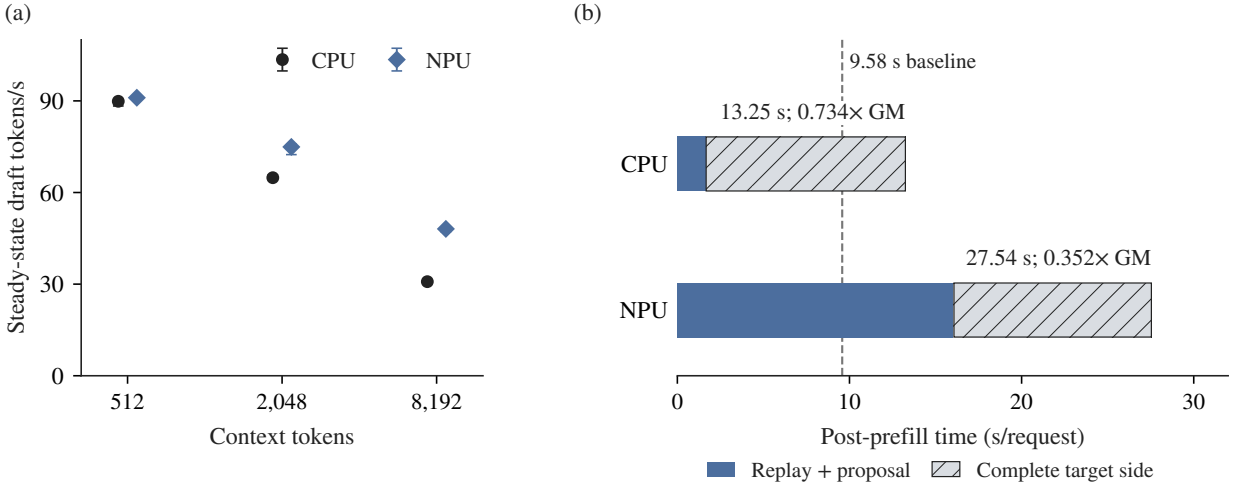
(a)
CPU
NPU
Steady-state draft tokens/s
90
60
30
0
512
2,048
8,192
Context tokens
(b)
9.58 s baseline
13.25 s; 0.734× GM
CPU
27.54 s; 0.352× GM
NPU
0
10
20
30
Post-prefill time (s/request)
Replay + proposal
Complete target side