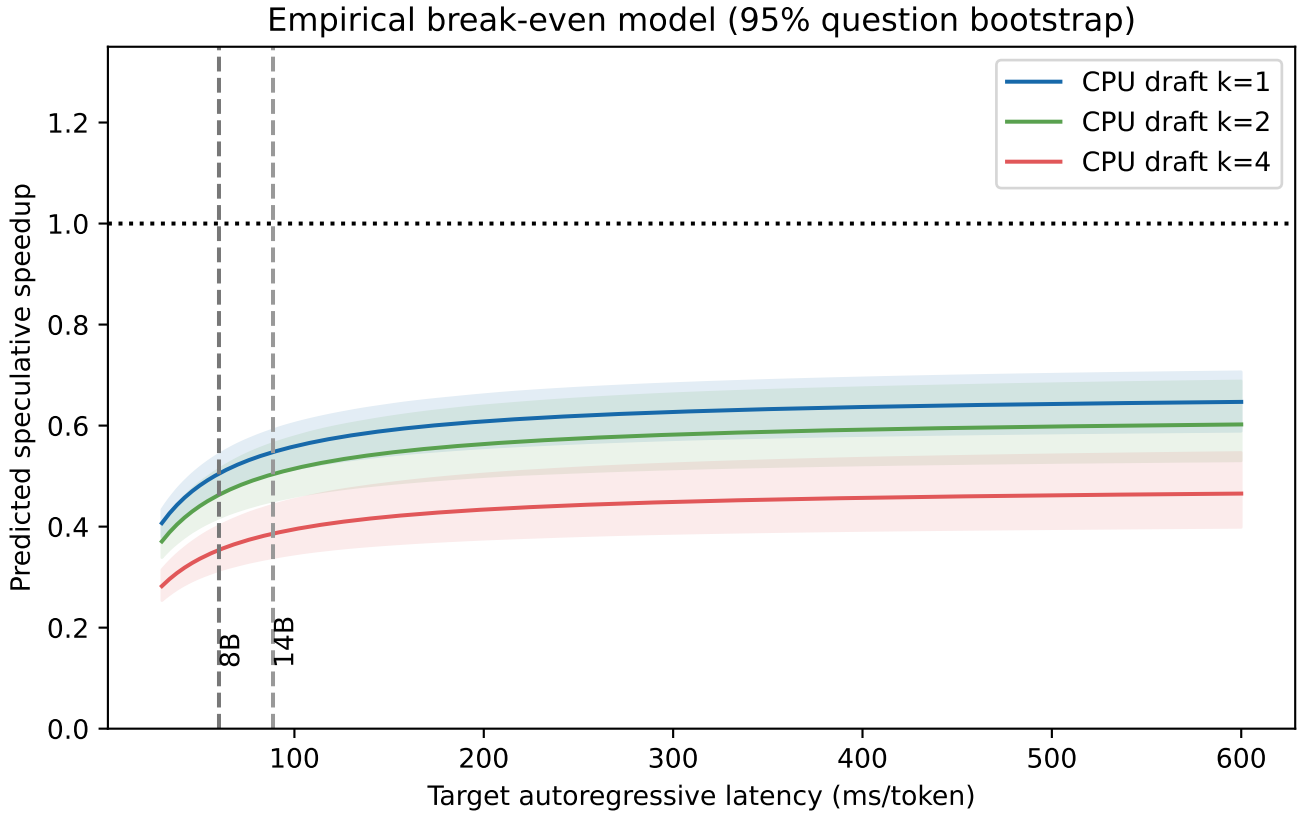
Empirical break-even model (95% question bootstrap)
CPU draft k=1
CPU draft k=2
CPU draft k=4
Predicted speculative speedup
1.2
1.0
0.8
0.6
0.4
0.2
0.0
8B
14B
100
200
300
400
500
600
Target autoregressive latency (ms/token)